

```
In [1]: using MLDatasets
using PyPlot
using Random, Statistics
using Flux
```

Rapid intro to supervised learning with neural nets I: from scratch

This notebook gives a rapid introduction to supervised learning with neural networks. The example is based on [Chapter 1 of Nielsen's online book "Neural Networks and Deep Learning"](#) and it guides you to set up the neural network training using the [Flux](#) Julia package.

For further reading I recommend also the review article ["A high-bias, low-variance introduction to Machine Learning for physicists"](#).

A few words about Flux

[Flux](#) provides functionality for deep learning that is similar to the established ML libraries [TensorFlow](#) and [PyTorch](#). This means, in particular, automatic differentiation and simple ways to construct deep learning models (=neural networks).


Automatic differentiation

With [automatic differentiation](#) we can let the computer compute gradients of arbitrary functions. In Flux the function `gradient()` takes a function and an input for the fuction as argument and the gradient of the given function. Example:

```
In [2]: f(x)=x^2

f_prime(x)=Tracker.data(Flux.gradient(f,x)[1])

x=-1.0:0.1:1.0
plot(x,f.(x),label=L"$f(x)$")
plot(x,f_prime.(x),label=L"$f'(x)$")
xlabel(L"$x$")
legend()
```



```
Out[2]: PyObject <matplotlib.legend.Legend object at 0x7fa1911fdf40>
```

Notice: For the purpose of automatic differentiation Flux introduces *tracked* data types. Therefore, the plain `Flux.gradient(f,x)[1]` returns a tracked `Float` and we have to call the `Tracker.data` function in order to convert it into a plain `Float`.

The MNIST hand-written digits data set

Let's first get a simple exemplary data set - the MNIST hand-written digits. The following cell downloads both the test and training parts of the data set.

```
In [3]: # load full training set
train_x, train_y = float.(MNIST.traindata())
# load full test set
test_x, test_y = float.(MNIST.testdata());
```

`trainData` is now a array of shape `(28,28, 60000)`, meaning that we have 60k images of 28×28 pixels (grayscale), each showing one hand-written digit. `trainLabels` holds the corresponding *labels*, i.e. an integer for each image, stating which digit it shows.

Defining a neural network model using Flux

```
In [4]: function build_model(layers; imgsize=(28,28))
m = Dense(prod(imgsize), layers[1], sigmoid)
for j in 2:length(layers)
m = Chain(m, Dense(layers[j-1], layers[j],sigmoid))
end
return m
end
```

```
Out[4]: build_model (generic function with 1 method)
```

Now we can again what the network thinks about our images of digits. For this purpose we define `initialize_network` and `neural_network` analogous to part I of the tutorial, but this time based on our `MyNet` class.

```
In [5]: net_layers=[100,10]

neural_network = build_model(net_layers)
params = Flux.params(neural_network)

neural_network(reshape(train_x[:, :, 1], 28*28))
```

```
Out[5]: Tracked 10-element Vector{Float32}:
 0.7877969f0
 0.56963897f0
 0.4330673f0
 0.59149516f0
 0.50833476f0
 0.6606239f0
 0.50392526f0
 0.48794314f0
 0.48309577f0
 0.5678873f0
```

Next, we need a cost function. This is the same as in the previous notebook.

```
In [6]: function cost_function(predictions, labels)
"""This function evaluates the cost function for given predictions and labels
Args:
* predictions: Predictions from neural net. Array of shape mathcal T x 10.
* labels: Correct labels for the corresponding images. Array of mathcal T integers.
Returns: Cost associated with the neural network predictions for the given data.
"""

labels = Flux.onehotbatch(labels, 0:9)

cost = sum((predictions-labels).^2)
return cost / size(labels)[2]
end
```

```
Out[6]: cost_function (generic function with 1 method)
```

With this, we can check the performance of our randomly initialized network in classifying some of our images:

```
In [7]: batch = train_x[:, :, 1:128] # select a batch of images
labels = train_y[1:128] # and corresponding labels

# ! compute neural network predictions
predictions = neural_network(reshape(batch, 28*28, size(batch)[3]))

# ! evaluate the cost function
cost_function(predictions, labels)
```

```
Out[7]: 3.0449834f0 (tracked)
```

Now, what is missing is a function to compute the gradients of the cost function. This is easily solved using `Flux.gradient()` for automatic differentiation:

```
In [8]: function cost_function_gradient(net, params, batch, labels)
return Flux.gradient(() -> cost_function(net(reshape(batch, 28*28, size(batch)[3])), labels), params)
end
```

```
Out[8]: cost_function_gradient (generic function with 1 method)
```

Finally, we are ready to train the network:

```
In [9]: function evaluate_predictions(predictions, labels)
"""This is a helper function that counts how many of the given predictions match the labels.
Args:
* `predictions`: Predictions from neural network (=activations on output layer)
* `labels`: correct labels
Returns: Number of correct predictions, i.e., number of cases, in which the index of the maximal
activation matches the given label.
"""
pred_labels = [Int(findmax(predictions[:,i])[2])-1 for i in 1:size(predictions)[2]]
return sum(pred_labels .== labels)
end

prng_key = Random.seed!(1234)

neural_network = build_model(net_layers)
params = Flux.params(neural_network)

# Here we define the hyperparamters
num_epochs = 10 # Number of epochs to loop over
learning_rate = 0.001 # Learning rate
batch_size = 128 # Size of mini-batches

# Compute the number of mini-batches that matches the chosen mini-batch size
batch_number = floor(Int, size(train_x)[end] / batch_size)

# Evaluate network and assess performance
predictions = neural_network(reshape(test_x, 28*28, size(test_x)[3]))
current_cost = cost_function(predictions, test_y)
correct_predictions = evaluate_predictions(predictions, test_y)
println("Initial cost: $(current_cost)")
println("Correctly predicted labels: $(correct_predictions) / $(length(test_y))")

for n in 1:num_epochs

println("Episode $(n)")
order = shuffle(1:length(train_y))
samples, labels = ( reshape(train_x[:, :, order][ :, :, 1:Int(batch_number*batch_size)], 28,28,128, :),
reshape(train_y[order][1:Int(batch_number*batch_size)], 128, :))

for i in 1:batch_number

# Compute gradients
gs=cost_function_gradient(neural_network, params, samples[:, :, i], labels[:, i])

# Perform SGD parameter update step
for p in params
Flux.Optimise.update!(p, -learning_rate*gs[p])
end

end

# Evaluate network and assess performance
predictions = neural_network(reshape(test_x, 28*28, size(test_x)[3]))
current_cost = cost_function(predictions, test_y)
correct_predictions = evaluate_predictions(predictions, test_y)
println("Current cost: $(current_cost)")
println("Correctly predicted labels: $(correct_predictions)/$(length(test_y))")
end
```

```
Initial cost: 2.9444165f0 (tracked)
Correctly predicted labels: 1067 / 10000
Episode 1
Current cost: 1.0339497f0 (tracked)
Correctly predicted labels: 2503/10000
Episode 2
Current cost: 0.9229258f0 (tracked)
Correctly predicted labels: 3013/10000
Episode 3
Current cost: 0.89791876f0 (tracked)
Correctly predicted labels: 3280/10000
Episode 4
Current cost: 0.8880436f0 (tracked)
Correctly predicted labels: 3401/10000
Episode 5
Current cost: 0.8825424f0 (tracked)
Correctly predicted labels: 3452/10000
Episode 6
Current cost: 0.8786173f0 (tracked)
Correctly predicted labels: 3505/10000
Episode 7
Current cost: 0.87532216f0 (tracked)
Correctly predicted labels: 3533/10000
Episode 8
Current cost: 0.8722753f0 (tracked)
Correctly predicted labels: 3553/10000
Episode 9
Current cost: 0.8693074f0 (tracked)
Correctly predicted labels: 3572/10000
Episode 10
Current cost: 0.8663381f0 (tracked)
Correctly predicted labels: 3612/10000
```